







✧ Hoshi: A* Pathfinding Algorithm in MIPS Assembly

Overview

Hoshi (星, Japanese for “star”) is a complete implementation of the A* pathfinding algorithm in MIPS assembly language. This project demonstrates low-level programming concepts by implementing a complex algorithmic solution in assembly language, offering valuable insights into computer architecture and optimization techniques.

Features

-  **Complete A* algorithm implementation** in pure MIPS assembly
-  **Priority Queue** with efficient insertion and extraction operations
-  **Manhattan/Chebyshev Distance Heuristic** for optimal path calculation
-  **Visual Representation** through bitmap display for step-by-step algorithm progression
-  **Obstacle Detection** with robust path planning around barriers
-  **Path Reconstruction** with visual highlighting of the optimal route

A* Algorithm: Core Concepts and Implementation

What is A*?

A* (pronounced “A-star”) is an informed search algorithm widely used in pathfinding and graph traversal. It efficiently plots a traversable path between multiple nodes by maintaining a priority queue of paths and choosing the lowest-cost path to expand.

Algorithm Overview

The A* algorithm is built on three key components: 1. **g-score**: The actual cost from the start node to the current node 2. **h-score**: A heuristic estimate of the cost from the current node to the goal 3. **f-score**: The sum of g-score and h-score, representing the total estimated cost

The algorithm maintains two sets: - **Open Set**: Nodes to be evaluated (stored in a priority queue) - **Closed Set**: Nodes already evaluated

Algorithm Pseudocode

1. Initialize:
 - Open List: Contains nodes yet to be evaluated (start with start node)
 - Closed List: Stores nodes already evaluated (starts empty)
 - Set $\text{start.g} = 0$: Cost from start to start
 - Set $\text{start.h} = \text{heuristic}(\text{start}, \text{goal})$: Estimated cost from start to goal
 - Compute $\text{start.f} = \text{start.g} + \text{start.h}$
 - Set $\text{start.parent} = \text{null}$: No parent yet

2. Main loop: While open list is not empty:
 - Select current as the node in openList with the lowest f value
 - Goal check:
 - If current == goal, return the reconstructed path
 - Move current node:
 - Remove current from openList
 - Add current to closedList
 - Process neighbors:
 - For each neighbor of current:
 - Skip if neighbor is in closedList
 - Compute tentative_g = current.g + distance(current, neighbor)
 - If neighbor is not in openList: Add it
 - Else if tentative_g >= neighbor.g: Skip (existing path is better)
 - Otherwise (this path is better):
 - Update neighbor.parent = current
 - Update neighbor.g = tentative_g
 - Update neighbor.h = heuristic(neighbor, goal)
 - Recompute neighbor.f = neighbor.g + neighbor.h
3. If loop ends with no path found:
 - Return failure: No path exists between start and goal
4. Path Reconstruction Function:
 - Start from goal node
 - Trace back using parent links, adding each node to a path
 - Return the path in reverse (from start to goal)

Core Implementation in MIPS Assembly

The A* algorithm is implemented in the `a_star` function, which follows the pseudocode above:

```
a_star:
    # Function setup
    addi    $sp, $sp, -4
    sw      $ra, 0($sp)

    # Initialize nodes and data structures
    jal     initialize_nodes

    # Setup start node
    lw      $s0, start_x
    lw      $s1, start_y

    # Set g-score of start node to 0
    move    $a0, $s0
    move    $a1, $s1
    li      $a2, 0
    jal     set_g_score
```

```

# Calculate h-score using heuristic
lw      $a2, goal_x
lw      $a3, goal_y
jal     chebyshevDistance # or manhattanDistance

# Calculate f-score (g + h)
move    $a2, $v0
move    $a0, $s0
move    $a1, $s1
jal     set_f_score

# Push start node to open set
move    $a0, $s0
move    $a1, $s1
li      $a2, 0 # parent (none for start)
move    $a3, $v0 # f-score
jal     push

# Main loop
a_star_loop:
# Check if open set is empty
la      $t0, heapSize
lw      $t0, 0($t0)
beqz    $t0, a_star_no_path

# Pop node with lowest f-score
jal     pop

# Check if popped node is goal
la      $t0, extracted_node
lw      $t1, 0($t0) # x
lw      $t2, 4($t0) # y
lw      $t3, goal_x
lw      $t4, goal_y

beq     $t1, $t3, goal_check_y
j       not_goal

goal_check_y:
beq     $t2, $t4, a_star_found_path

not_goal:
# Add current node to closed set
move    $a0, $t1
move    $a1, $t2
jal     add_to_closed_set

```

```

    # Process neighbors (4 directions)
    li      $s0, 0 # direction counter

process_neighbors:
    li      $t0, 4
    beq     $s0, $t0, a_star_loop # If all neighbors processed

    # Calculate neighbor coordinates
    la      $t0, d4x
    la      $t1, d4y
    sll     $t2, $s0, 2
    add     $t0, $t0, $t2
    add     $t1, $t1, $t2
    lw      $t2, 0($t0) # dx
    lw      $t3, 0($t1) # dy

    la      $t0, extracted_node
    lw      $t4, 0($t0) # current.x
    lw      $t5, 4($t0) # current.y

    add     $t6, $t4, $t2 # neighbor.x
    add     $t7, $t5, $t3 # neighbor.y

    # Check if valid position
    move    $a0, $t6
    move    $a1, $t7
    jal     is_valid_position
    beqz    $v0, next_neighbor

    # Check if in closed set
    move    $a0, $t6
    move    $a1, $t7
    jal     is_in_closed_set
    bnez    $v0, next_neighbor

    # Calculate tentative g-score
    move    $a0, $t4
    move    $a1, $t5
    jal     get_g_score

    addi    $t0, $v0, 1 # tentative_g = current.g + 1

    # Get neighbor's current g-score
    move    $a0, $t6
    move    $a1, $t7
    jal     get_g_score

    # Compare tentative_g with neighbor.g

```

```

bge      $t0, $v0, next_neighbor # Skip if not better

# Update neighbor's scores
move     $a0, $t6
move     $a1, $t7
move     $a2, $t0 # new g-score
jal      set_g_score

# Calculate h-score
move     $a0, $t6
move     $a1, $t7
lw       $a2, goal_x
lw       $a3, goal_y
jal      chebyshevDistance

# Calculate new f-score
add      $a2, $t0, $v0 # g + h
move     $a0, $t6
move     $a1, $t7
jal      set_f_score

# Set parent
move     $a0, $t6
move     $a1, $t7
move     $a2, $t4 # parent.x
move     $a3, $t5 # parent.y
jal      set_parent

# Push neighbor to open set
move     $a0, $t6
move     $a1, $t7
move     $a2, $a2 # parent reference
move     $a3, $a2 # f-score
jal      push

next_neighbor:
    addi   $s0, $s0, 1
    j      process_neighbors

a_star_found_path:
    # Print success message
    la     $a0, path_found_msg
    li     $v0, 4
    syscall

# Reconstruct path
lw       $a0, goal_x
lw       $a1, goal_y
jal      constructPathProcedure

```

Data Structures and Modules

The implementation is organized into several modules, each handling specific aspects of the algorithm:

1. Priority Queue

The priority queue is implemented as a binary min-heap, which ensures efficient extraction of the node with the lowest f-score.

Node Structure in the Priority Queue

Offset	Field	Size (bytes)
0	x	4
4	y	4
8	parent	4
12	fScore	4

Key Operations:

- **push:** Inserts a node with $O(\log n)$ complexity
- **pop:** Extracts the node with lowest f-score with $O(\log n)$ complexity

2. Node List

Each node in the grid has specific properties that track its state in the A* algorithm.

Node Structure

Offset	Field	Size (bytes)
0	x	4
4	y	4
8	wall	4
12	gScore	4
16	hScore	4
20	fScore	4
24	parent_x	4
28	parent_y	4

Key Operations:

- **initialize_nodes:** Sets up the grid based on map data
- **set_g_score/get_g_score:** Manages cost from start
- **set_f_score/get_f_score:** Manages total estimated cost

3. Bitmap Display

The bitmap module manages visualization, providing a graphical representation of the A* algorithm's execution.

Display Constants

```
.eqv    displayWidth, 16      # Width of the display in units
.eqv    displayHeight, 16     # Height of the display in units
.eqv    gridCellWidth, 2      # Width of each grid cell
.eqv    gridCellHeight, 2     # Height of each grid cell
.eqv    gridWidth, 8          # Width of the grid in cells
.eqv    gridHeight, 8         # Height of the grid in cells
.eqv    bitmapBaseAddress, 0x10040000 # Memory address of bitmap
```

Key Operations:

- **clearScreen**: Initializes the display
- **drawGridNode**: Renders a single node with specific color based on its state
- **drawGrid**: Renders the entire grid

4. Heuristic Functions

The A* algorithm uses heuristic functions to estimate the cost from any node to the goal.

Available Heuristics:

- **Manhattan Distance**: Sum of horizontal and vertical distances
- **Chebyshev Distance**: Maximum of horizontal and vertical distances

manhattanDistance:

```
# Calculate |x1-x2| + |y1-y2|
sub    $t0, $a0, $a2          # x1-x2
abs    $t0, $t0               # |x1-x2|

sub    $t1, $a1, $a3          # y1-y2
abs    $t1, $t1               # |y1-y2|

add    $v0, $t0, $t1          # |x1-x2| + |y1-y2|
jr     $ra
```

chebyshevDistance:

```
# Calculate max(|x1-x2|, |y1-y2|)
sub    $t0, $a0, $a2          # x1-x2
abs    $t0, $t0               # |x1-x2|

sub    $t1, $a1, $a3          # y1-y2
abs    $t1, $t1               # |y1-y2|

# Find maximum
bge    $t0, $t1, max_is_x
move   $v0, $t1
j      chebyshev_return
```

max_is_x:

```
move   $v0, $t0
```

5. Path Reconstruction

Once the A* algorithm finds a path, it traces back from the goal to the start using parent pointers.

Implementation Details

Memory Management

The implementation uses a consistent pattern to locate nodes in memory:

```
# Calculate node address from (x,y) coordinates
lw      $t0, map_width      # Load grid width
mul      $t1, $a1, $t0      # t1 = y * width
add      $t1, $t1, $a0      # t1 = y * width + x
lw      $t2, node_size      # Load node size in bytes
mul      $t3, $t1, $t2      # t3 = index * node_size
la      $t4, nodes          # Load base address
add      $t4, $t4, $t3      # t4 = base + offset
```

This efficiently implements the formula: $\&\text{nodes}[y * \text{width} + x]$ to convert 2D coordinates to memory addresses.

Register Usage Strategy

The implementation follows a consistent register allocation strategy: - $\$s0-\$s7$: Preserved across function calls, used for loop variables and important data - $\$t0-\$t9$: Temporary calculations, not preserved across calls - $\$a0-\$a3$: Function arguments - $\$v0-\$v1$: Function return values - $\$ra$: Return address register, preserved on stack when making nested calls

Stack Management

Proper stack management is critical for function calls and recursion:

```
# Function prologue
addi     $sp, $sp, -4      # Allocate stack space
sw       $ra, 0($sp)      # Save return address

# Function body
# ...

# Function epilogue
lw       $ra, 0($sp)      # Restore return address
addi     $sp, $sp, 4      # Deallocate stack space
jr       $ra              # Return
```


Visualization

The visualization uses a consistent color scheme: - **Color 0** (White): Background/free space - **Color 1** (Black): Walls/obstacles - **Color 2** (Green): Goal node - **Color 3** (Green): Final path - **Color 5** (Yellow): Current node being explored - **Color 8** (Cyan): Start node - **Color 9** (Gray): Nodes in the open set

This color coding makes it easy to understand the algorithm's progress visually.

Optimization Techniques

1. **Priority Queue:** $O(\log n)$ operations for managing the open set
2. **Register Usage:** Strategic register allocation minimizes memory access
3. **Closed Set:** Efficient tracking of evaluated nodes
4. **Memory Access:** Calculates addresses efficiently to minimize overhead
5. **Data Organization:** Node structures organized for efficient access patterns

Performance Analysis

Operation	Time Complexity
Node Extraction	$O(\log n)$
Node Insertion	$O(\log n)$
Path Reconstruction	$O(p)$ where p is path length
Heuristic Calculation	$O(1)$
Overall Algorithm	$O(E \log V)$ where V is number of nodes and E is number of edges

Movement Directions

The implementation supports both 4-way and 8-way movement:

```
# 4-way movement (up, right, down, left)
d4x:          .word      0, 1, 0, -1
d4y:          .word     -1, 0, 1, 0

# 8-way movement (includes diagonals)
d8x:          .word      0, 1, 1, 1, 0, -1, -1, -1
d8y:          .word     -1, -1, 0, 1, 1, 1, 0, -1
```

The default implementation uses 4-way movement for simplicity and clarity.

Educational Value

This implementation offers several educational insights: 1. **Low-level Programming:** Direct memory management and register allocation 2. **Algorithm Implementation:** From theory to practical assembly code 3. **Data Structures:** Priority queue and grid management 4.

Visualization: Real-time algorithm execution display 5. **Optimization:** Balancing readability with performance

HOSHI : the story of a program to reach a star that no one can see