







✧ Hoshi: A* Pathfinding Algorithm in MIPS Assembly

Overview

Hoshi (星, Japanese for “star”) is a complete implementation of the A* pathfinding algorithm in MIPS assembly language. This project demonstrates low-level programming concepts by implementing a complex algorithmic solution in assembly language, offering valuable insights into computer architecture and optimization techniques.

Features

-  **Complete A* algorithm implementation** in pure MIPS assembly
-  **Priority Queue** with efficient insertion and extraction operations
-  **Manhattan/Chebyshev Distance Heuristic** for optimal path calculation
-  **Visual Representation** through bitmap display for step-by-step algorithm progression
-  **Obstacle Detection** with robust path planning around barriers
-  **Path Reconstruction** with visual highlighting of the optimal route

A* Algorithm: Core Concepts and Implementation

What is A*?

A* (pronounced “A-star”) is an informed search algorithm widely used in pathfinding and graph traversal. It efficiently plots a traversable path between multiple nodes by maintaining a priority queue of paths and choosing the lowest-cost path to expand.

Algorithm Overview

The A* algorithm is built on three key components: 1. **g-score**: The actual cost from the start node to the current node 2. **h-score**: A heuristic estimate of the cost from the current node to the goal 3. **f-score**: The sum of g-score and h-score, representing the total estimated cost

The algorithm maintains two sets: - **Open Set**: Nodes to be evaluated (stored in a priority queue) - **Closed Set**: Nodes already evaluated

Algorithm Pseudocode

1. Initialize:
 - Open List: Contains nodes yet to be evaluated (start with start node)
 - Closed List: Stores nodes already evaluated (starts empty)
 - Set $\text{start.g} = 0$: Cost from start to start
 - Set $\text{start.h} = \text{heuristic}(\text{start}, \text{goal})$: Estimated cost from start to goal
 - Compute $\text{start.f} = \text{start.g} + \text{start.h}$
 - Set $\text{start.parent} = \text{null}$: No parent yet

2. Main loop: While open list is not empty:
 - Select current as the node in openList with the lowest f value
 - Goal check:
 - If current == goal, return the reconstructed path
 - Move current node:
 - Remove current from openList
 - Add current to closedList
 - Process neighbors:
 - For each neighbor of current:
 - Skip if neighbor is in closedList
 - Compute tentative_g = current.g + distance(current, neighbor)
 - If neighbor is not in openList: Add it
 - Else if tentative_g >= neighbor.g: Skip (existing path is better)
 - Otherwise (this path is better):
 - Update neighbor.parent = current
 - Update neighbor.g = tentative_g
 - Update neighbor.h = heuristic(neighbor, goal)
 - Recompute neighbor.f = neighbor.g + neighbor.h
3. If loop ends with no path found:
 - Return failure: No path exists between start and goal
4. Path Reconstruction Function:
 - Start from goal node
 - Trace back using parent links, adding each node to a path
 - Return the path in reverse (from start to goal)

Core psudo Implementation

The A* algorithm is implemented in the a_star function, which follows the pseudocode above:

```
function A_Star(start, goal): // Initialize open and closed lists
  openList = [start] // Nodes to be evaluated
  closedList = [] // Nodes already evaluated
```

```
  // Initialize node properties
  start.g = 0 // Cost from start to start is 0
  start.h = heuristic(start, goal) // Estimate to goal
  start.f = start.g + start.h // Total estimated cost
  start.parent = null // For path reconstruction
  while openList is not empty:
    // Get node with lowest f value - implement using a priority queue
    // for faster retrieval of the best node
    current = node in openList with lowest f value

    // Check if we've reached the goal
    if current == goal:
      return reconstruct_path(current)

    // Move current node from open to closed list
```

```

remove current from openList
add current to closedList

// Check all neighboring nodes
for each neighbor of current:
    if neighbor in closedList:
        continue // Skip already evaluated nodes

    // Calculate tentative g score
    tentative_g = current.g + distance(current, neighbor)

    if neighbor not in openList:
        add neighbor to openList
    else if tentative_g >= neighbor.g:
        continue // This path is not better

    // This path is the best so far
    neighbor.parent = current
    neighbor.g = tentative_g
    neighbor.h = heuristic(neighbor, goal)
    neighbor.f = neighbor.g + neighbor.h

return failure // No path exists

function reconstruct_path(current): path = [] while current is not null: add current to beginning
of path current = current.parent return path

```

Data Structures and Modules

The implementation is organized into several modules, each handling specific aspects of the algorithm:

1. Priority Queue

The priority queue is implemented as a binary min-heap, which ensures efficient extraction of the node with the lowest f-score.

Node Structure in the Priority Queue

Offset	Field	Size (bytes)
0	x	4
4	y	4
8	parent	4
12	fScore	4

Key Operations:

- **push:** Inserts a node with $O(\log n)$ complexity
- **pop:** Extracts the node with lowest f-score with $O(\log n)$ complexity

2. Node List

Each node in the grid has specific properties that track its state in the A* algorithm.

Node Structure

Offset	Field	Size (bytes)
0	x	4
4	y	4
8	wall	4
12	gScore	4
16	hScore	4
20	fScore	4
24	parent_x	4
28	parent_y	4

Key Operations:

- **initialize_nodes:** Sets up the grid based on map data
- **set_g_score/get_g_score:** Manages cost from start
- **set_f_score/get_f_score:** Manages total estimated cost

3. Bitmap Display

The bitmap module manages visualization, providing a graphical representation of the A* algorithm's execution.

Display Constants

```
.eqv    displayWidth, 16      # Width of the display in units
.eqv    displayHeight, 16    # Height of the display in units
.eqv    gridCellWidth, 2     # Width of each grid cell
.eqv    gridCellHeight, 2    # Height of each grid cell
.eqv    gridWidth, 8         # Width of the grid in cells
.eqv    gridHeight, 8        # Height of the grid in cells
.eqv    bitmapBaseAddress, 0x10040000 # Memory address of bitmap
```

Key Operations:

- **clearScreen:** Initializes the display
- **drawGridNode:** Renders a single node with specific color based on its state
- **drawGrid:** Renders the entire grid

4. Heuristic Functions

The A* algorithm uses heuristic functions to estimate the cost from any node to the goal.

Available Heuristics:

- **Manhattan Distance:** Sum of horizontal and vertical distances
- **Chebyshev Distance:** Maximum of horizontal and vertical distances

```

manhattanDistance:
    # Calculate |x1-x2| + |y1-y2|
    sub    $t0, $a0, $a2    # x1-x2
    abs    $t0, $t0        # |x1-x2|

    sub    $t1, $a1, $a3    # y1-y2
    abs    $t1, $t1        # |y1-y2|

    add    $v0, $t0, $t1    # |x1-x2| + |y1-y2|
    jr     $ra

```

```

chebyshevDistance:
    # Calculate max(|x1-x2|, |y1-y2|)
    sub    $t0, $a0, $a2    # x1-x2
    abs    $t0, $t0        # |x1-x2|

    sub    $t1, $a1, $a3    # y1-y2
    abs    $t1, $t1        # |y1-y2|

    # Find maximum
    bge    $t0, $t1, max_is_x
    move    $v0, $t1
    j      chebyshev_return

```

```

max_is_x:
    move    $v0, $t0

```

```

chebyshev_return:
    jr     $ra

```

5. Path Reconstruction

Once the A* algorithm finds a path, it traces back from the goal to the start using parent pointers.

Implementation Details

Memory Management

The implementation uses a consistent pattern to locate nodes in memory:

```

# Calculate node address from (x,y) coordinates
lw      $t0, map_width    # Load grid width
mul      $t1, $a1, $t0    # t1 = y * width
add      $t1, $t1, $a0    # t1 = y * width + x
lw      $t2, node_size    # Load node size in bytes
mul      $t3, $t1, $t2    # t3 = index * node_size
la      $t4, nodes        # Load base address
add      $t4, $t4, $t3    # t4 = base + offset

```

This efficiently implements the formula: $\&\text{nodes}[y * \text{width} + x]$ to convert 2D coordinates to memory addresses.

Register Usage Strategy

The implementation follows a consistent register allocation strategy: - \$s0-\$s7: Preserved across function calls, used for loop variables and important data - \$t0-\$t9: Temporary calculations, not preserved across calls - \$a0-\$a3: Function arguments - \$v0-\$v1: Function return values - \$ra: Return address register, preserved on stack when making nested calls

Stack Management

Proper stack management is critical for function calls and recursion:

```
# Function prologue
addi    $sp, $sp, -4      # Allocate stack space
sw      $ra, 0($sp)      # Save return address

# Function body
# ...

# Function epilogue
lw      $ra, 0($sp)      # Restore return address
addi    $sp, $sp, 4      # Deallocate stack space
jr      $ra              # Return
```

Visualization

The visualization uses a consistent color scheme: - **Color 0** (White): Background/free space - **Color 1** (Black): Walls/obstacles - **Color 2** (Green): Goal node - **Color 3** (Green): Final path - **Color 5** (Yellow): Current node being explored - **Color 8** (Cyan): Start node - **Color 9** (Gray): Nodes in the open set

This color coding makes it easy to understand the algorithm’s progress visually.

Optimization Techniques

- 1. **Priority Queue:** $O(\log n)$ operations for managing the open set
- 2. **Register Usage:** Strategic register allocation minimizes memory access
- 3. **Closed Set:** Efficient tracking of evaluated nodes
- 4. **Memory Access:** Calculates addresses efficiently to minimize overhead
- 5. **Data Organization:** Node structures organized for efficient access patterns

Performance Analysis

Operation	Time Complexity
Node Extraction	$O(\log n)$

Operation	Time Complexity
Node Insertion	$O(\log n)$
Path Reconstruction	$O(p)$ where p is path length
Heuristic Calculation	$O(1)$
Overall Algorithm	$O(E \log V)$ where V is number of nodes and E is number of edges

Movement Directions

The implementation supports both 4-way and 8-way movement:

4-way movement (up, right, down, left)

d4x: .word 0, 1, 0, -1

d4y: .word -1, 0, 1, 0

8-way movement (includes diagonals)

d8x: .word 0, 1, 1, 1, 0, -1, -1, -1

d8y: .word -1, -1, 0, 1, 1, 1, 0, -1

The default implementation uses 4-way movement for simplicity and clarity.

Educational Value

This implementation offers several educational insights: 1. **Low-level Programming:** Direct memory management and register allocation 2. **Algorithm Implementation:** From theory to practical assembly code 3. **Data Structures:** Priority queue and grid management 4.

Visualization: Real-time algorithm execution display 5. **Optimization:** Balancing readability with performance

